



Level 0x09

Kernel Interfacing



Topics

- Events
- Hacker History
- Linux syscalls


Upcoming Events

- December hacking contests
 - Advent of Cyber - [TryHackMe.com](https://tryhackme.com)
 - Sans Institute [Holiday Hack Challenge](#) - Less CTF, more game, probably more incident responder and defender focused
 - [Advent of Code](#) - Programming challenges. Used to be 25 2-parters, now reducing down to 12 days of challenges
 - [Pwn.College](#) is going to do an Advent of Pwn



Richard Stallman aka rms

- Harvard Physics graduate
- Worked at MIT AI Lab... where a printer triggered him...
- Started the GNU Project (GNUs Not Unix)
 - GNU Compiler Collection (GCC), GNU make
 - GNU Debugger (GDB)
 - GPL License (Open Source License)
 - Free Software Foundation (Freedom: Libre vs Gratis)
- MIT AI Lab had leadership involved with Jeffrey Epstein
- Has an insanely large speech rider
 - What pets are preferred at the place of stay
 - What codecs you are allowed to distribute speech in
 - Don't even mention breakfast



Advent of Pwn Day 1

- Hint tells you to run objdump
 - -d is for disassemble

What we find in checklist:


- Some stuff
- syscall
- subb / addb instructions
 - Sooo many of these
 - Really, pages of them
- cmpb and jne instructions

```
Terminal - ubuntu@2025~day-01: /challenge
File Edit View Terminal Tabs Help
./check-list:      file format elf64-x86-64

Disassembly of section .text:
0000000000401000 <.text>:
401000:   48 89 e5          mov    %rsp,%rbp
401003:   48 81 ec 00 05 00 00 sub   $0x500,%rsp
40100a:   b8 00 00 00 00       mov    $0x0,%eax
40100f:   bf 00 00 00 00       mov    $0x0,%edi
401014:   48 8d b5 00 fc ff ff lea   -0x400(%rbp),%rsi
40101b:   ba 00 04 00 00       mov    $0x400,%edx
401020:   0f 05              syscall
401022:   80 6d b7 c0          subb  $0xc0,-0x49(%rbp)
401026:   80 85 35 fe ff ff 0a addb  $0xa,-0x1cb(%rbp)
40102d:   80 85 30 fe ff ff b7 addb  $0xb7,-0x1d0(%rbp)
401034:   80 ad 04 fe ff ff 60 subb  $0x60,-0x1fc(%rbp)
40103b:   80 6d a8 c3          subb  $0xc3,-0x58(%rbp)
40103f:   80 6d b5 63          subb  $0x63,-0x4b(%rbp)
401043:   80 6d fd c5          subb  $0xc5,-0x3(%rbp)
401047:   80 85 2a fc ff ff 4e addb  $0x4e,-0x3d6(%rbp)
40104e:   80 85 e6 fe ff ff a7 addb  $0xa7,-0x11a(%rbp)
:
aa1438:   80 85 2f fe ff ff ed addb  $0xed,-0x1d1(%rbp)
aa143f:   80 85 36 fc ff ff c4 addb  $0xc4,-0x3ca(%rbp)
aa1446:   80 ad 0c ff ff ff d7 subb  $0xd7,-0xf4(%rbp)
aa144d:   80 ad b6 fc ff ff 2a subb  $0x2a,-0x34a(%rbp)
aa1454:   80 bd 00 fc ff ff 2b cmpb  $0x2b,-0x400(%rbp)
aa145b:   0f 85 09 33 00 00       jne   0xaa476a
aa1461:   80 bd 01 fc ff ff 8f cmpb  $0x8f,-0x3ff(%rbp)
aa1468:   0f 85 fc 32 00 00       jne   0xaa476a
aa146e:   80 bd 02 fc ff ff 22 cmpb  $0x22,-0x3fe(%rbp)
aa1475:   0f 85 ef 32 00 00       jne   0xaa476a
aa147b:   80 bd 03 fc ff ff af cmpb  $0xaf,-0x3fd(%rbp)
aa1482:   0f 85 e2 32 00 00       jne   0xaa476a
aa1488:   80 bd 04 fc ff ff c7 cmpb  $0xc7,-0x3fc(%rbp)
aa148f:   0f 85 d5 32 00 00       jne   0xaa476a
aa1495:   80 bd 05 fc ff ff 19 cmpb  $0x19,-0x3fb(%rbp)
aa149c:   0f 85 c8 32 00 00       jne   0xaa476a
aa14a2:   80 bd 06 fc ff ff bc cmpb  $0xbc,-0x3fa(%rbp)
aa14a9:   0f 85 bb 32 00 00       jne   0xaa476a
aa14af:   80 bd 07 fc ff ff 74 cmpb  $0x74,-0x3f9(%rbp)
aa14b6:   0f 85 ae 32 00 00       jne   0xaa476a
aa14bc:   80 bd 08 fc ff ff 46 cmpb  $0x46,-0x3f8(%rbp)
```

What is inside an OS kernel

- Hardware Interfacing
 - Memory Access
 - SSD / Hard Drive control
 - Graphics Drivers
 - Keyboard / mouse interfacing
 - Ethernet / Wi-Fi
- Process control (threads)
- Filesystems
- Memory access
- Networking
- Enforces privilege




Trap / Kernel syscall

- We program / use applications in userspace
 - Code written in Java, Python, C/C++
 - Code is converted to machine code for our CPU
- We call libraries (other functions) to help us...
 - Read / write data from user or terminal
 - Draw images on the screen
 - GUI interfaces
- Libraries talk to kernel to do all of the above
- Kernel interfaces to the hardware


Example of reading data from a file

Trap (syscall) is special instruction, tells CPU to change privilege level and run some code as the kernel



System Call Table

- How does the kernel know what we want it to do we say “kernel do stuff” / syscall?



System Call Table

- How does the kernel know what we want it to do we say “kernel do stuff” / syscall?
- Pass arguments to kernel via CPU registers
 - RAX register is the sys call number

%rax	Name	Manual	Entry point
0	read	read(2)	sys_read
1	write	write(2)	sys_write
2	open	open(2)	sys_open
3	close	close(2)	sys_close
4	stat	stat(2)	sys_newstat
5	fstat	fstat(2)	sys_newfstat




Figure 1-17. The 10 steps in making the system call `read(fd, buffer, nbytes)`.

The syscall table is CPU arch specific. Many are similar, but each arch have differences


x86_64 (64-bit)

Compiled from [Linux 4.14.0 headers](#).

NR	syscall name	references	%rax	arg0 (%rdi)	arg1 (%rsi)	arg2 (%rdx)	arg3 (%r10)	arg4 (%r8)	arg5 (%r9)
0	read	man/ cs/	0x00	unsigned int fd	char *buf	size_t count	-	-	-
1	write	man/ cs/	0x01	unsigned int fd	const char *buf	size_t count	-	-	-
2	open	man/ cs/	0x02	const char *filename	int flags	umode_t mode	-	-	-
3	close	man/ cs/	0x03	unsigned int fd	-	-	-	-	-
4	stat	man/ cs/	0x04	const char *filename	struct __old_kernel_stat *statbuf	-	-	-	-
5	fstat	man/ cs/	0x05	unsigned int fd	struct __old_kernel_stat *statbuf	-	-	-	-
6	lstat	man/ cs/	0x06	const char *filename	struct __old_kernel_stat *statbuf	-	-	-	-
7	poll	man/ cs/	0x07	struct pollfd *ufds	unsigned int nfds	int timeout	-	-	-
8	lseek	man/ cs/	0x08	unsigned int fd	off_t offset	unsigned int whence	-	-	-
9	mmap	man/ cs/	0x09	?	?	?	?	?	?
10	mprotect	man/ cs/	0x0a	unsigned long start	size_t len	unsigned long prot	-	-	-
11	munmap	man/ cs/	0x0b	unsigned long addr	size_t len	-	-	-	-
12	brk	man/ cs/	0x0c	unsigned long brk	-	-	-	-	-

Linux vs Windows

- Linux
 - 300-400 syscalls
 - They rarely ever get changed
- WinNT kernel
 - 400+ syscalls
 - Change a lot / don't use directly!
 - Closed source / not documented
- Win32 API / ntdll.dll
 - 1000s of functions
 - How developers should access the kernel



Back to Advent of Pwn...

```
Terminal - ubuntu@2025~day-01: /challenge
File Edit View Terminal Tabs Help
./check-list:      file format elf64-x86-64

Disassembly of section .text:
0000000000401000 <.text>:
401000:   48 89 e5
401003:   48 81 ec 00 05 00 00
40100a:   b8 00 00 00 00
40100f:   bf 00 00 00 00
401014:   48 8d b5 00 fc ff ff
40101b:   ba 00 04 00 00
401020:   0f 05
401022:   80 6d b7 c0
401026:   80 85 35 fe ff ff 0a
40102d:   80 85 30 fe ff ff b7
401034:   80 ad 04 fe ff ff 60
40103b:   80 6d a8 c3
40103f:   80 6d b5 63
401043:   80 6d fd c5
401047:   80 85 2a fc ff ff 4e
40104e:   80 85 e6 fe ff ff a7
    mov    %rsp,%rbp
    sub    $0x500,%rsp
    mov    $0x0,%eax
    mov    $0x0,%edi
    lea    -0x400(%rbp),%rsi
    mov    $0x400,%edx
    syscall
    subb  $0xc0,-0x49(%rbp)
    addb  $0xa,-0x1cb(%rbp)
    addb  $0xb7,-0x1d0(%rbp)
    subb  $0x60,-0x1fc(%rbp)
    subb  $0xc3,-0x58(%rbp)
    subb  $0x63,-0x4b(%rbp)
    subb  $0xc5,-0x3(%rbp)
    addb  $0x4e,-0x3d6(%rbp)
    addb  $0xa7,-0x11a(%rbp)
```

x86_64 (64-bit)									
Compiled from Linux 4.14.0 headers.									
NR	syscall name	references	%rax	arg0 (%rdi)	arg1 (%rsi)	arg2 (%rdx)	arg3 (%r10)	arg4 (%r8)	arg5 (%r9)
0	read	man/cs/	0x00	unsigned int fd	char *buf	size_t count	-	-	-
1	write	man/cs/	0x01	unsigned int fd	const char *buf	size_t count	-	-	-
2	open	man/cs/	0x02	const char *filename	int flags	umode_t mode	-	-	-
3	close	man/cs/	0x03	unsigned int fd	-	-	-	-	-



Links

- Tanenbaum and Bos, Modern Operating Systems (book)
- Hausenblas, Learning Modern Linux (book)
- <https://filippo.io/linux-syscall-table/>
- <https://chromium.googlesource.com/chromiumos/docs/+/master/constants/syscalls.md>
- <https://j00ru.vexillium.org/syscalls/win32k/64/>
-